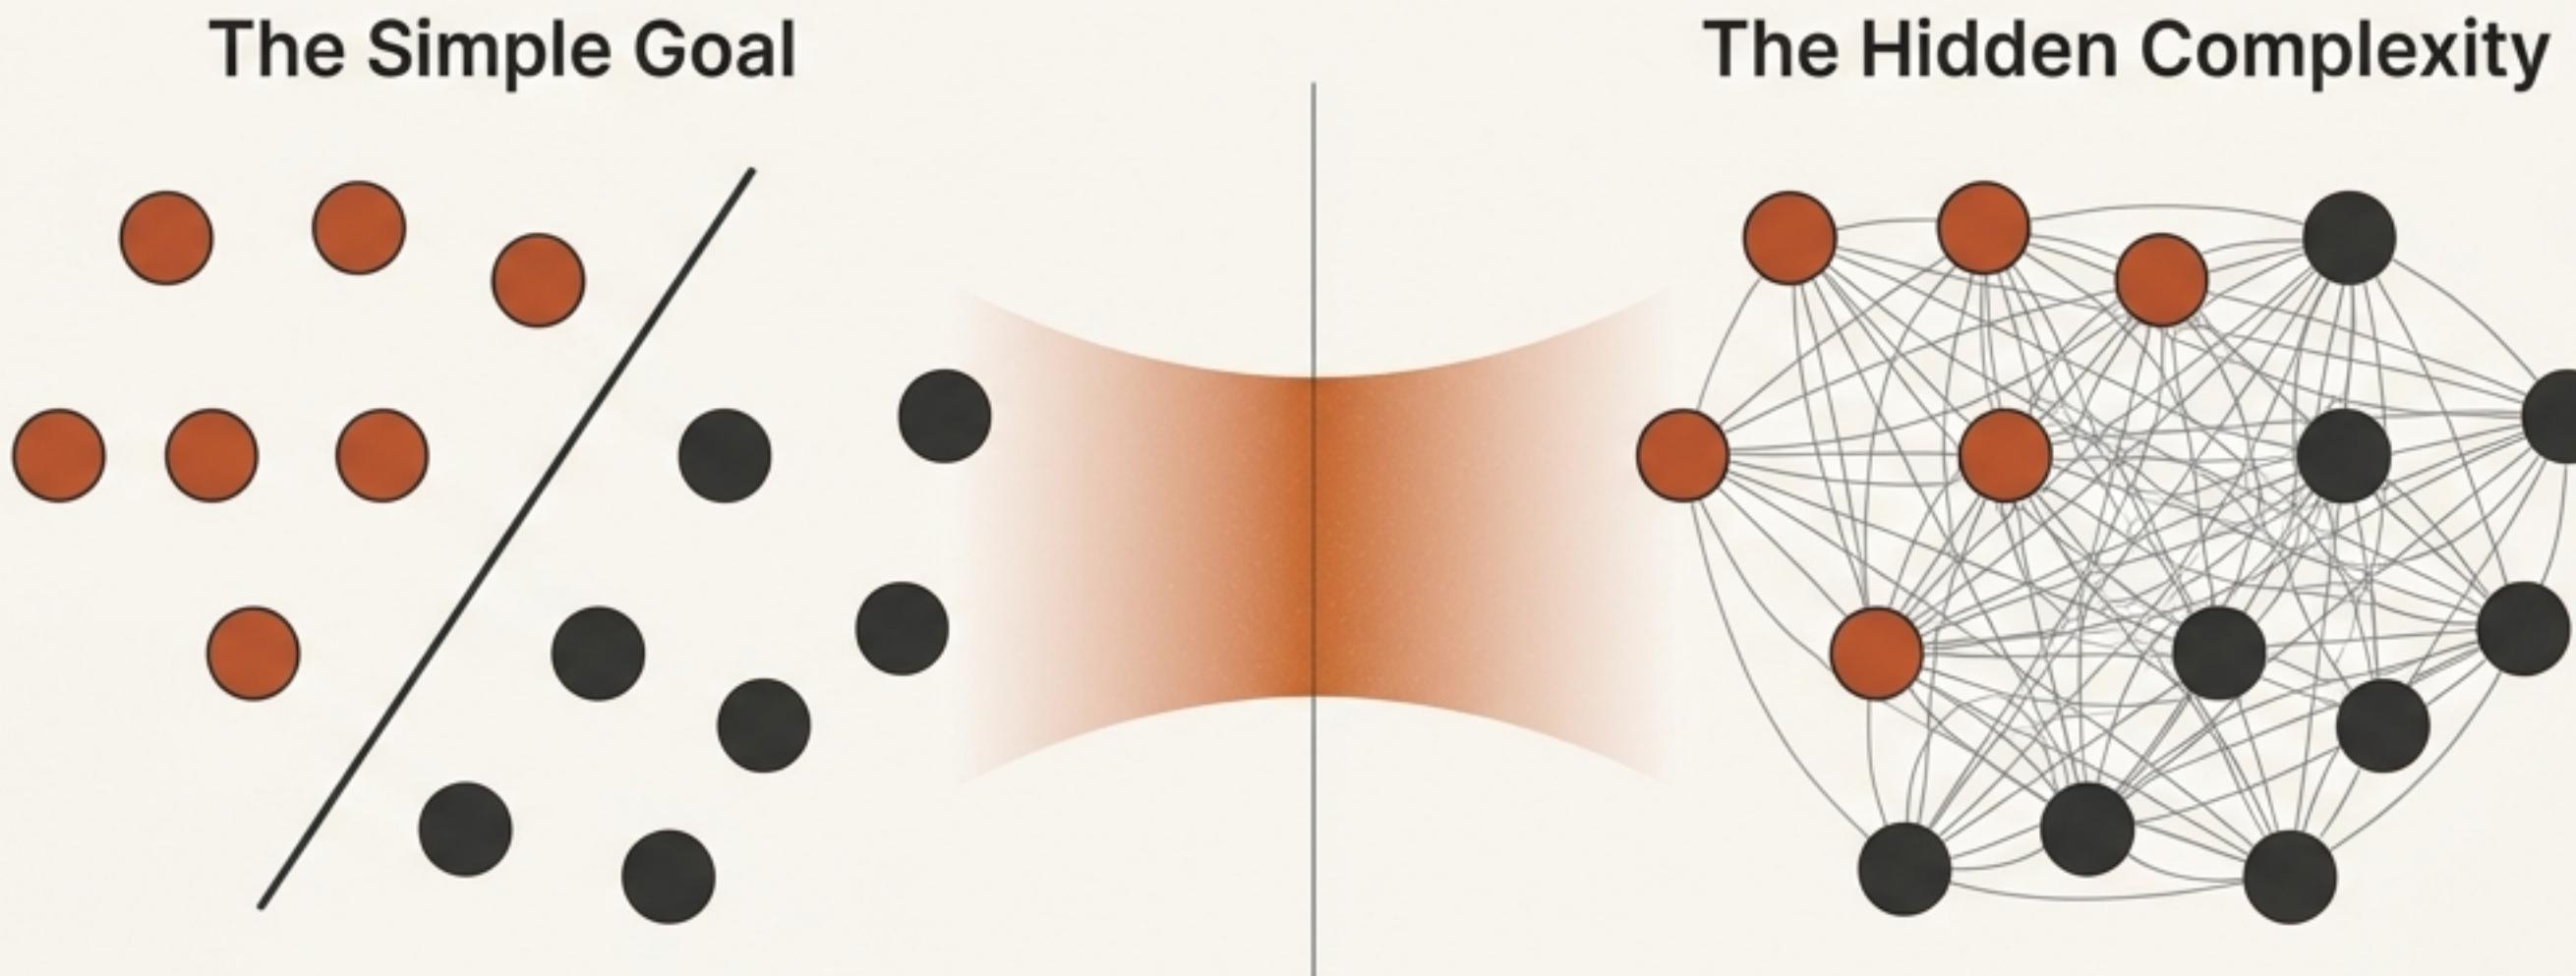


The Market Split Problem: A Benchmark's Journey from Logistics to Quantum Frontiers



An analytical explainer on the problem that broke the world's best solvers
and forced a paradigm shift in optimization.



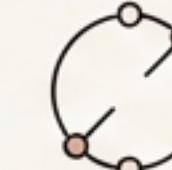
Practical
Origins



Legendary
Challenge



Algorithmic
Breakthrough



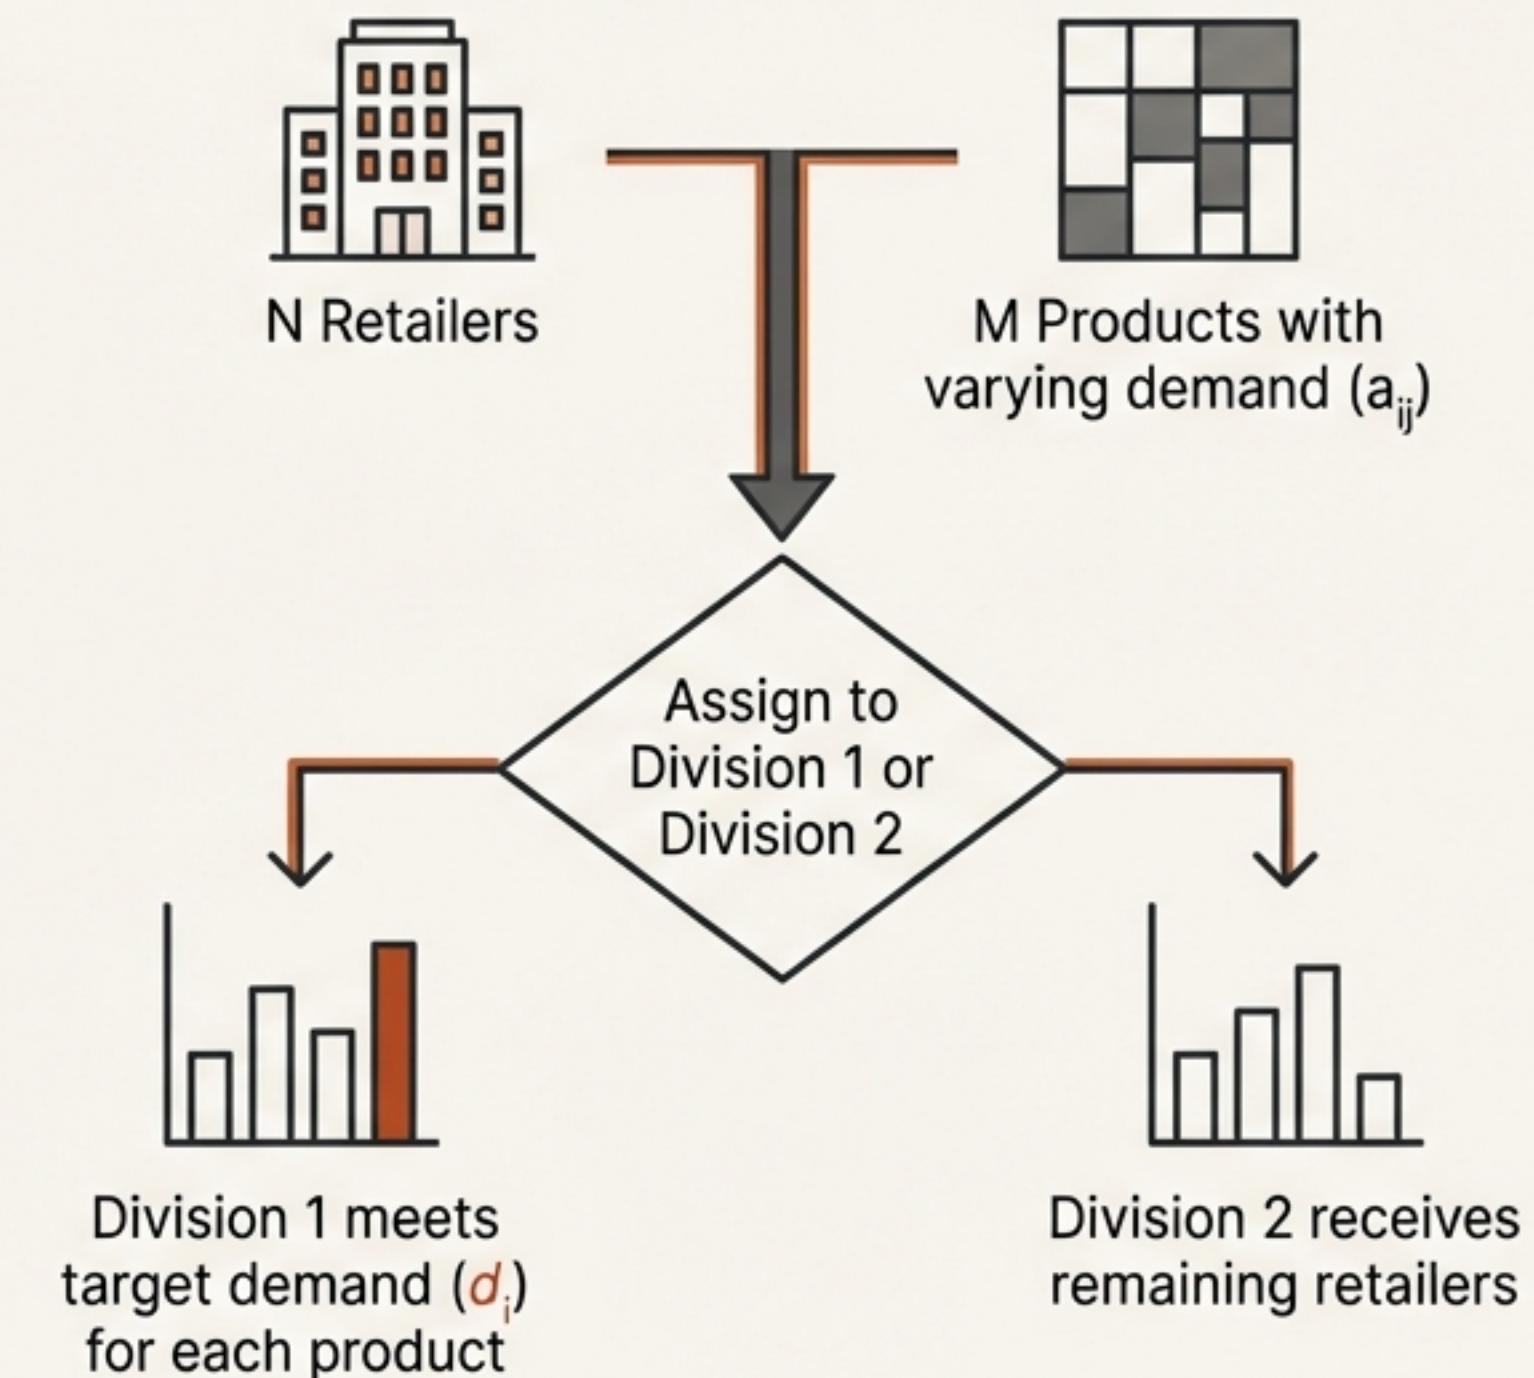
The New Quantum
Benchmark

It Began with a Practical Question in the 1970s UK Oil Market

The problem was first formally documented by H. Paul Williams in 1978 in 'Model Building in Mathematical Programming.'

- **The Scenario:** A company needed to allocate its retailers to two divisions (D_1, D_2) to balance market share across multiple products.
- **The Goal:** For each of m products, assign n retailers so that Division 1 controls a specific fraction (σ_i) of the total market demand (d_i).

For twenty years, it was considered a challenging but standard integer programming task, reflecting the messy, multi-constrained reality of industrial logistics.



The Problem is Mathematically Defined as a System of Linear Diophantine Equations

The Feasibility Variant (fMSP)

Find a binary vector $x \in \{0, 1\}^n$ that satisfies the system.

$$\sum_j (a_{ij} * x_j) = d_i, \text{ for } i = 1, \dots, m$$

Demand Binary Assignment Target

Complexity: NP-complete. The proof is by reduction from the classic Subset Sum Problem (SSP), which is the case when $m=1$. For $m > 1$, it becomes a significantly harder multi-dimensional subset sum problem.

The Optimization Variant (OPT)

For cases where a perfect split is impossible, minimize the total deviation from the targets.

$$\begin{aligned} & \min \sum_i |s_i| \\ \text{subject to } & \sum_j (a_{ij} * x_j) + s_i = d_i \\ & x_j \in \{0, 1\}, s_i \in \mathbb{Z} \end{aligned}$$

Slack/Deviation Variable

Note: This version is much harder for traditional solvers, as they must explore nearly the entire search tree to prove optimality.

In 1998, the Problem Was Weaponized into an Unsolvable Challenge

At the IPCO VI conference, Gérard Cornuéjols and Milind Dawande presented a set of instances that were virtually unsolvable for the world's leading commercial solvers.

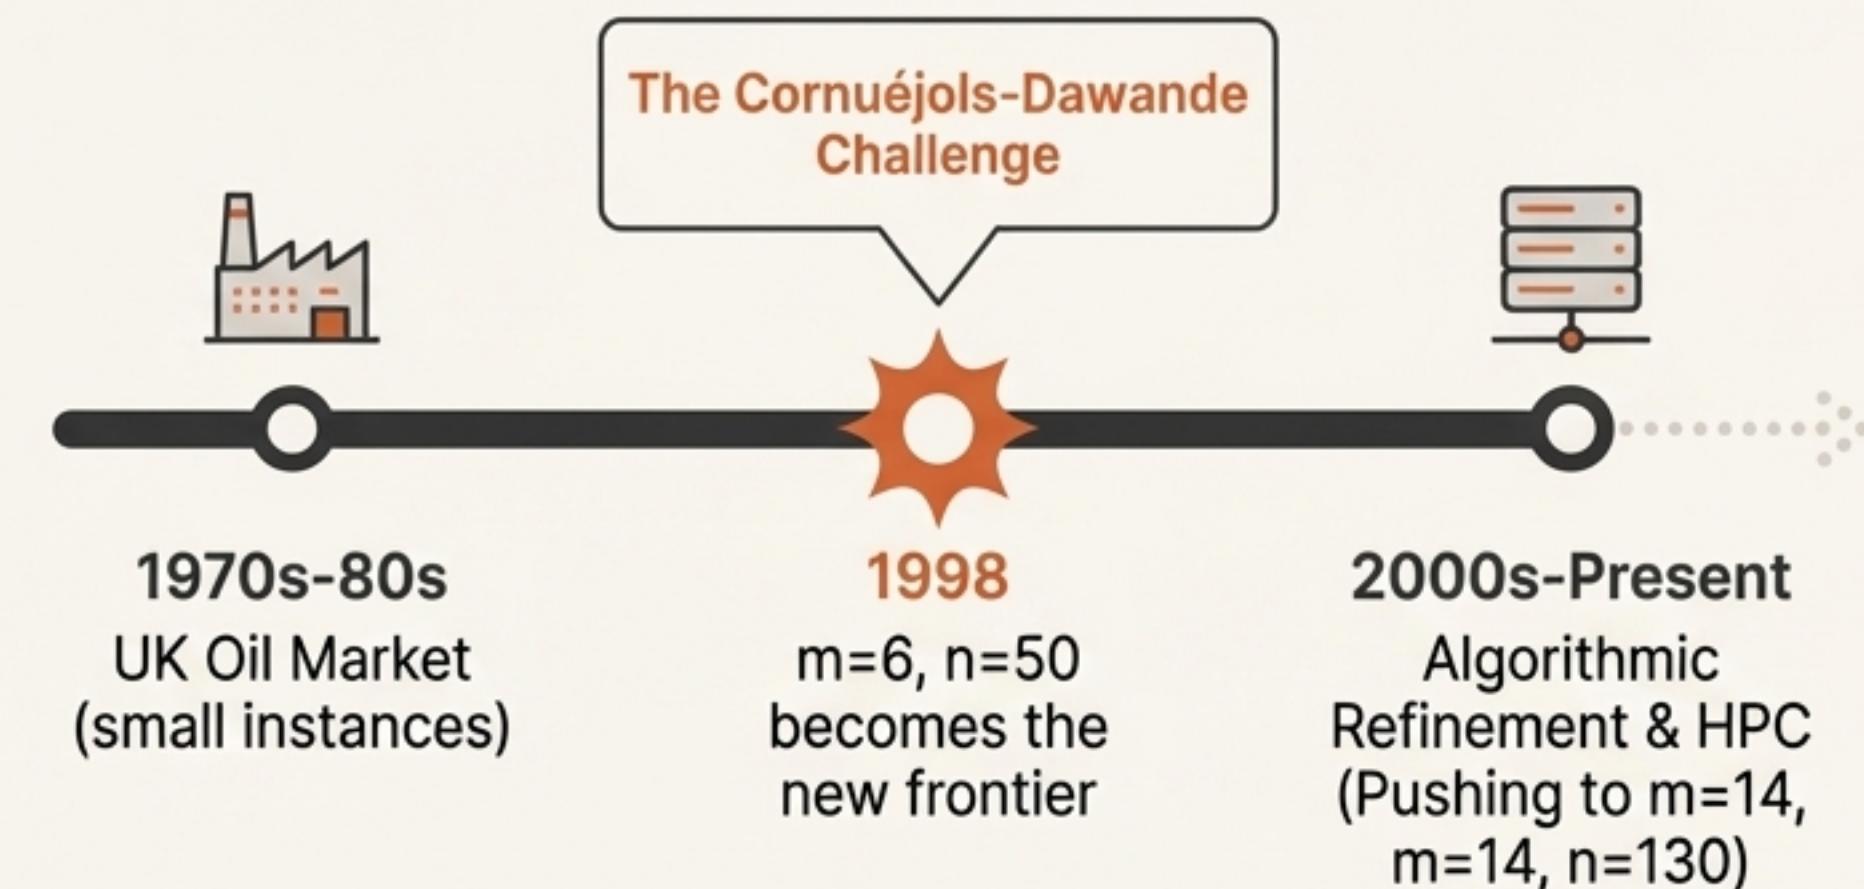
The "Hard" Instance Formula

They generated instances using the formula $n = 10(m-1)$, setting the number of variables to be roughly ten times the number of constraints.

Why It Worked

This formula created a search space that was:

- Sparsely populated with feasible integer solutions.
- Densely packed with near-feasible fractional points (e.g., solutions where many variables are 0.5).

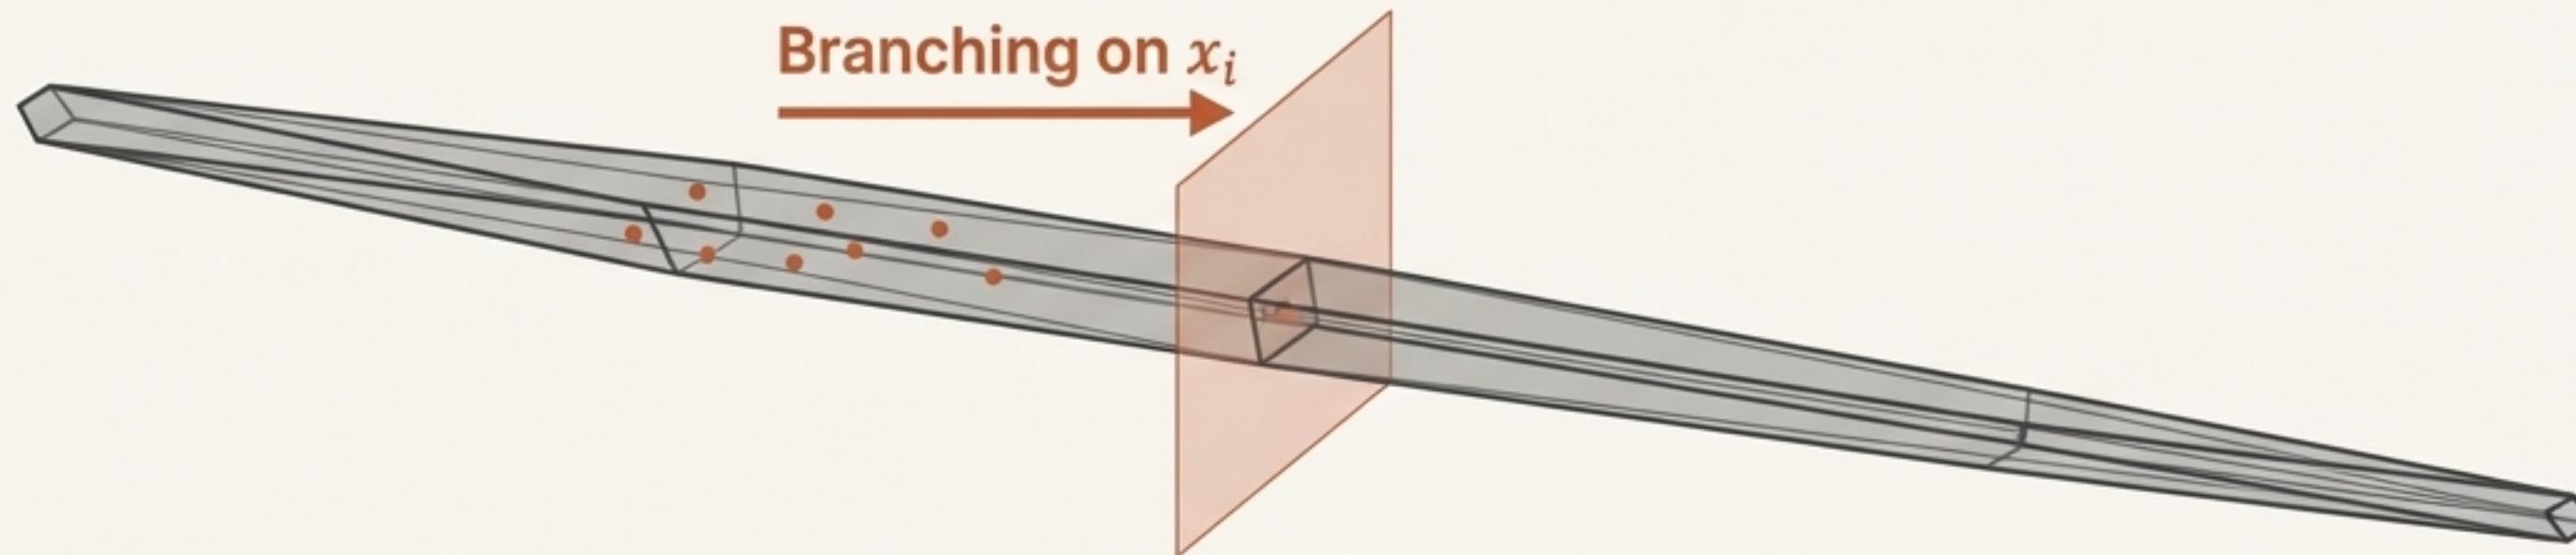


The Impact

This transformed the MIP into a benchmark that highlighted a critical gap in the capabilities of linear programming-based branch-and-bound methods.

Traditional Solvers Failed Due to the “Geometric Thinness” of the Feasible Region

The feasible region is a collection of discrete points within a highly constrained polytope.
When relaxed to $0 \leq x_j \leq 1$, the resulting LP polytope is extremely “thin”.



Problem 1: Ineffective Branching

Branching on a variable x_j cuts this thin polytope, but because it's so narrow in the coordinate directions, the search space is barely reduced.

The LP relaxation value often remains at zero for most nodes, failing to provide the lower bounds needed to prune subproblems, leading to an exponential explosion of nodes.

Problem 2: Symmetry and Fractional Solutions

The LP relaxation tends to find a “perfectly balanced” fractional solution where many variables are 0.5.

This provides no guidance toward an integer solution, forcing the solver into what is essentially an exhaustive brute-force search.

The Breakthrough Came from a Paradigm Shift: The Geometry of Numbers

The Old Way

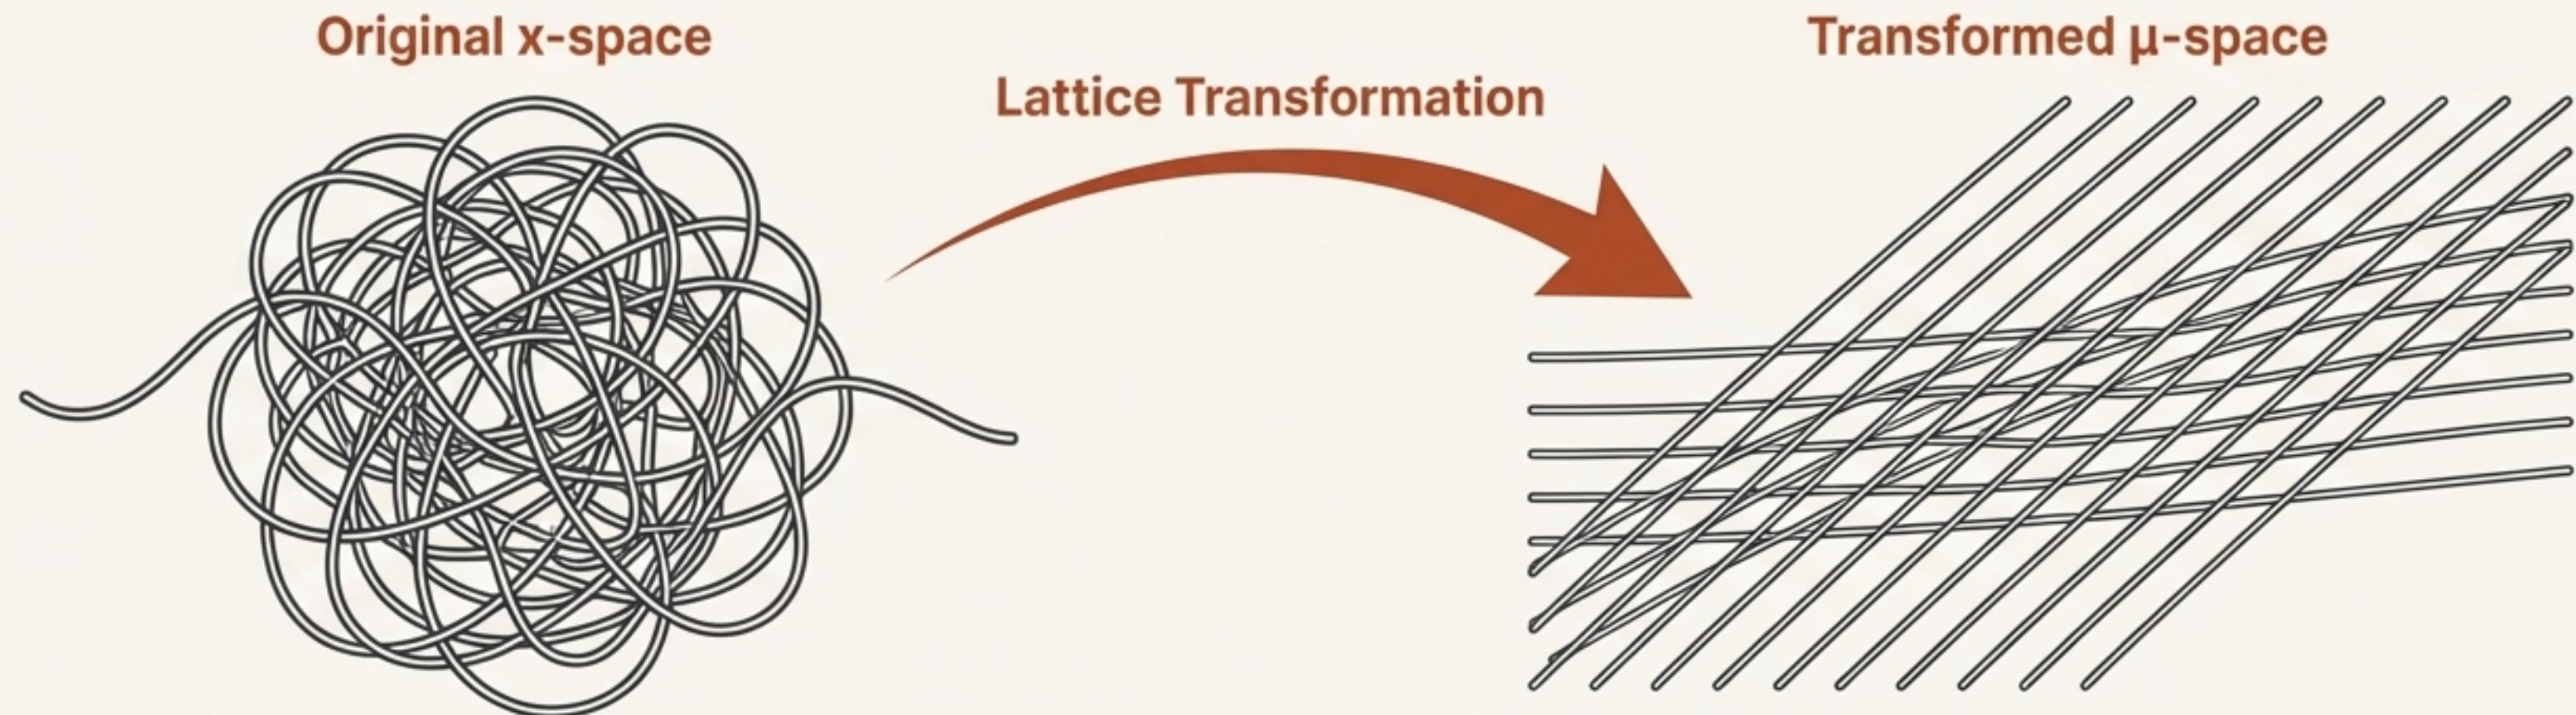
Treating the problem as a set of linear inequalities and searching for points within a thin polytope. This is the realm of Branch-and-Bound.

The New Way

Treating the problem as a search for a specific vector within a lattice. This is the realm of Lattice Basis Reduction.

Pioneers

Aardal, Hurkens, and Lenstra (AHL) developed the first successful lattice-based algorithm for the MSP.

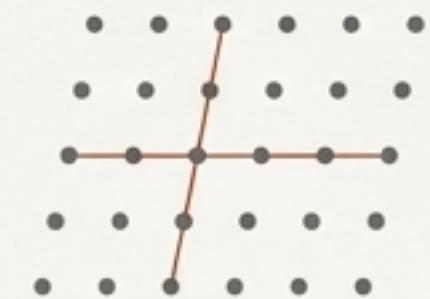


Central Idea: Instead of working in the original variable space, transform the problem into a new coordinate system where the constraints are easier to handle and the geometry is more favorable.

The AHL Algorithm Reformulates the Problem in a New Coordinate System

Step 1: Define the Lattice

The algorithm first considers the lattice of integer vectors that satisfy the homogeneous system $\mathbf{A}\mathbf{x} = \mathbf{0}$.



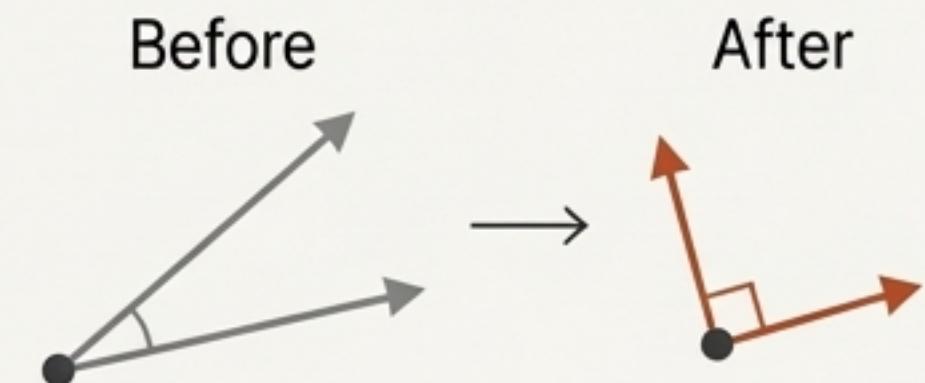
Step 2: Create the Extended Matrix

An extended matrix \mathbf{L} is constructed, incorporating the constraints (\mathbf{A}), the target vector (\mathbf{d}), and a large penalty parameter (\mathbf{N}).

$$\mathbf{L} = \begin{bmatrix} N * \mathbf{A} & -N * \mathbf{d} \\ \mathbf{0} & \mathbf{I} \end{bmatrix}$$

Step 3: Reduce the Basis

A lattice basis reduction algorithm (like LLL or BKZ) is applied to the columns of \mathbf{L} . This produces a new basis \mathbf{B} for the kernel of \mathbf{A} composed of 'short' and 'near-orthogonal' vectors, along with a particular solution $\hat{\mathbf{x}}$.



Key Insight: This process transforms the problem from finding \mathbf{x} to finding a vector of integer coefficients $\boldsymbol{\mu}$.

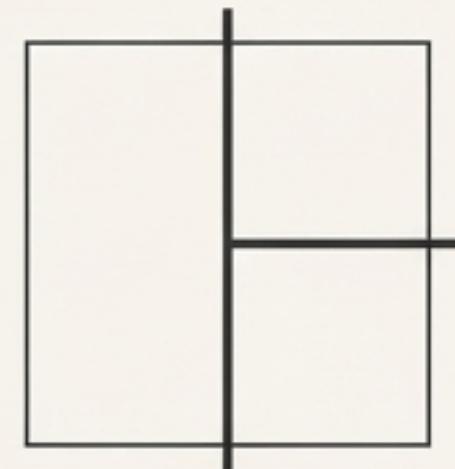
The Solution is Found by Searching in the Transformed "μ-space"

$$\mathbf{x} = \hat{\mathbf{x}} + \mathbf{B}\boldsymbol{\mu}$$

where $\boldsymbol{\mu}$ is a vector of integer coefficients in the new coordinate system.

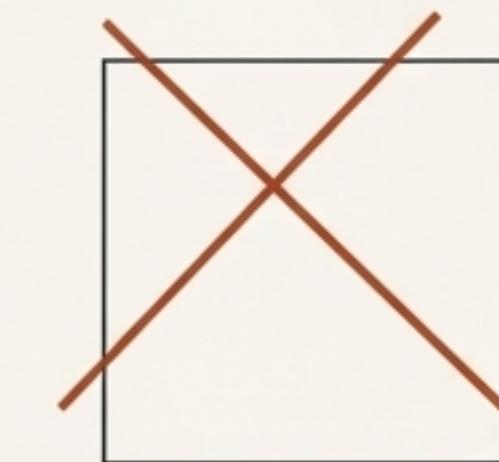
Because the new basis vectors in \mathbf{B} are very short, the range of possible integer values for the components of $\boldsymbol{\mu}$ is extremely small (often just 0 or 1 for the first few components).

Branching on x variables



Inefficient axis-aligned cuts that barely reduce the search volume in a thin polytope.

Branching on $\boldsymbol{\mu}$ variables (hyperplanes)



The New Strategy: The solver now branches on the $\boldsymbol{\mu}$ variables. Each $\boldsymbol{\mu}$ represents a linear combination of \mathbf{x} variables, allowing for powerful cuts that solve integer programs in polynomial time for a fixed number of variables.

The Performance Data Confirmed a Decisive Victory for Lattice Reformulation

A stark side-by-side look at the performance of standard branch-and-bound versus the new lattice-based approach on the Cornuéjols-Dawande challenge instances.

Instance Size (m, n)	Standard Branch-and-Bound Nodes	Lattice Reformulation Nodes	Success Rate
4, 30	> 1,000,000	< 500	100%
5, 40	Fail (Time Limit)	< 5,000	100%
6, 50	Fail (Memory/Time)	< 50,000	95%
7, 60	Unsolvable	≈ 200,000	90%

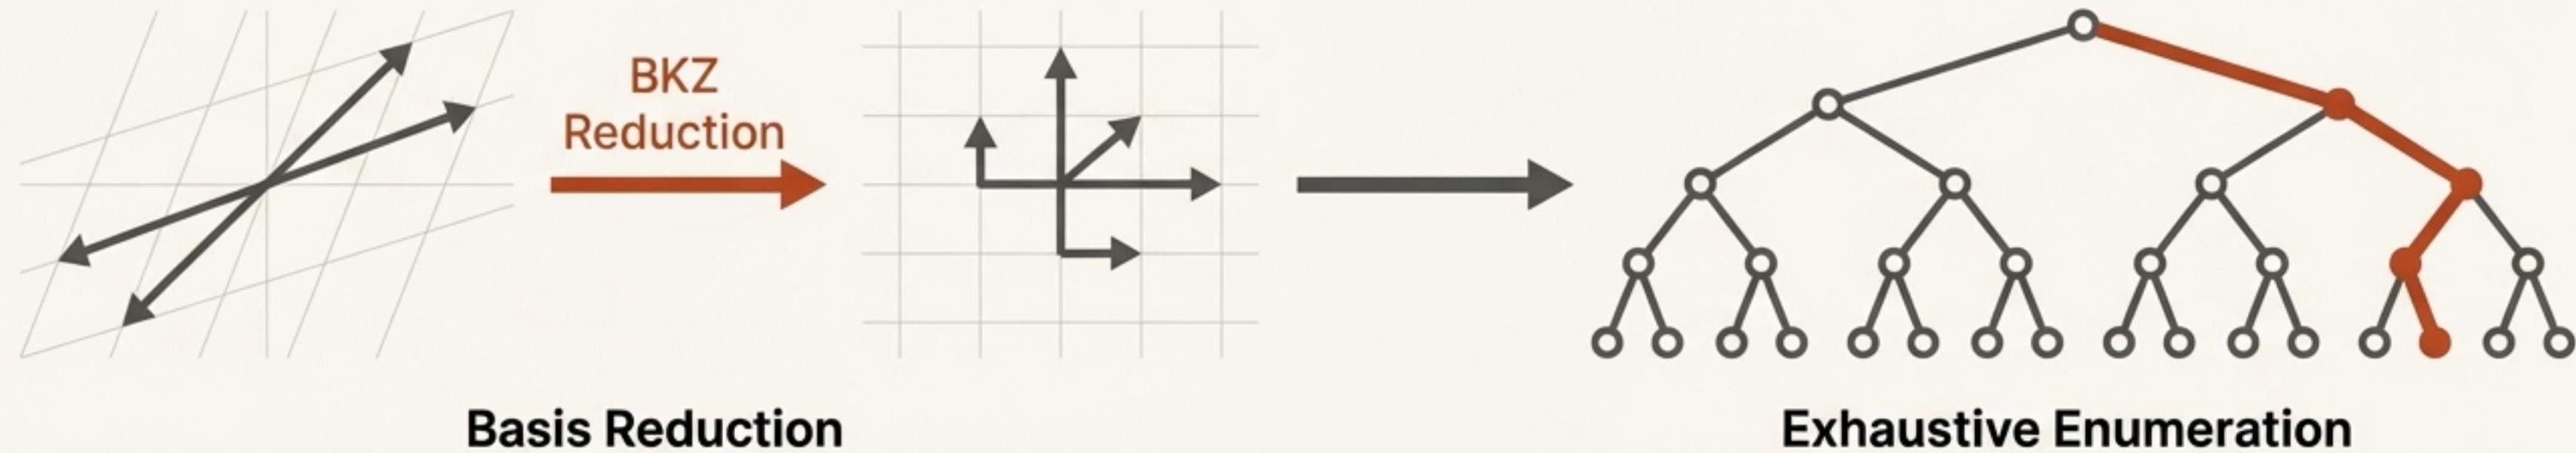
Source Note: Data based on seminal studies by Aardal et al. and Cornuéjols & Dawande (1998/1999).

The lattice approach didn't just improve performance; it solved problems that were previously considered completely out of reach.

Modern Solvers Refine Lattice Methods to Push the Solvability Frontier

State-of-the-Art Tool: `solvediophant`

Developed by Alfred Wassermann, this package represents a leading implementation of lattice-based methods.



Basis Reduction

Employs not just LLL but the more powerful (and computationally expensive) Blockwise Korkine-Zolotarev (BKZ) reduction to find higher-quality, shorter bases.

Exhaustive Enumeration

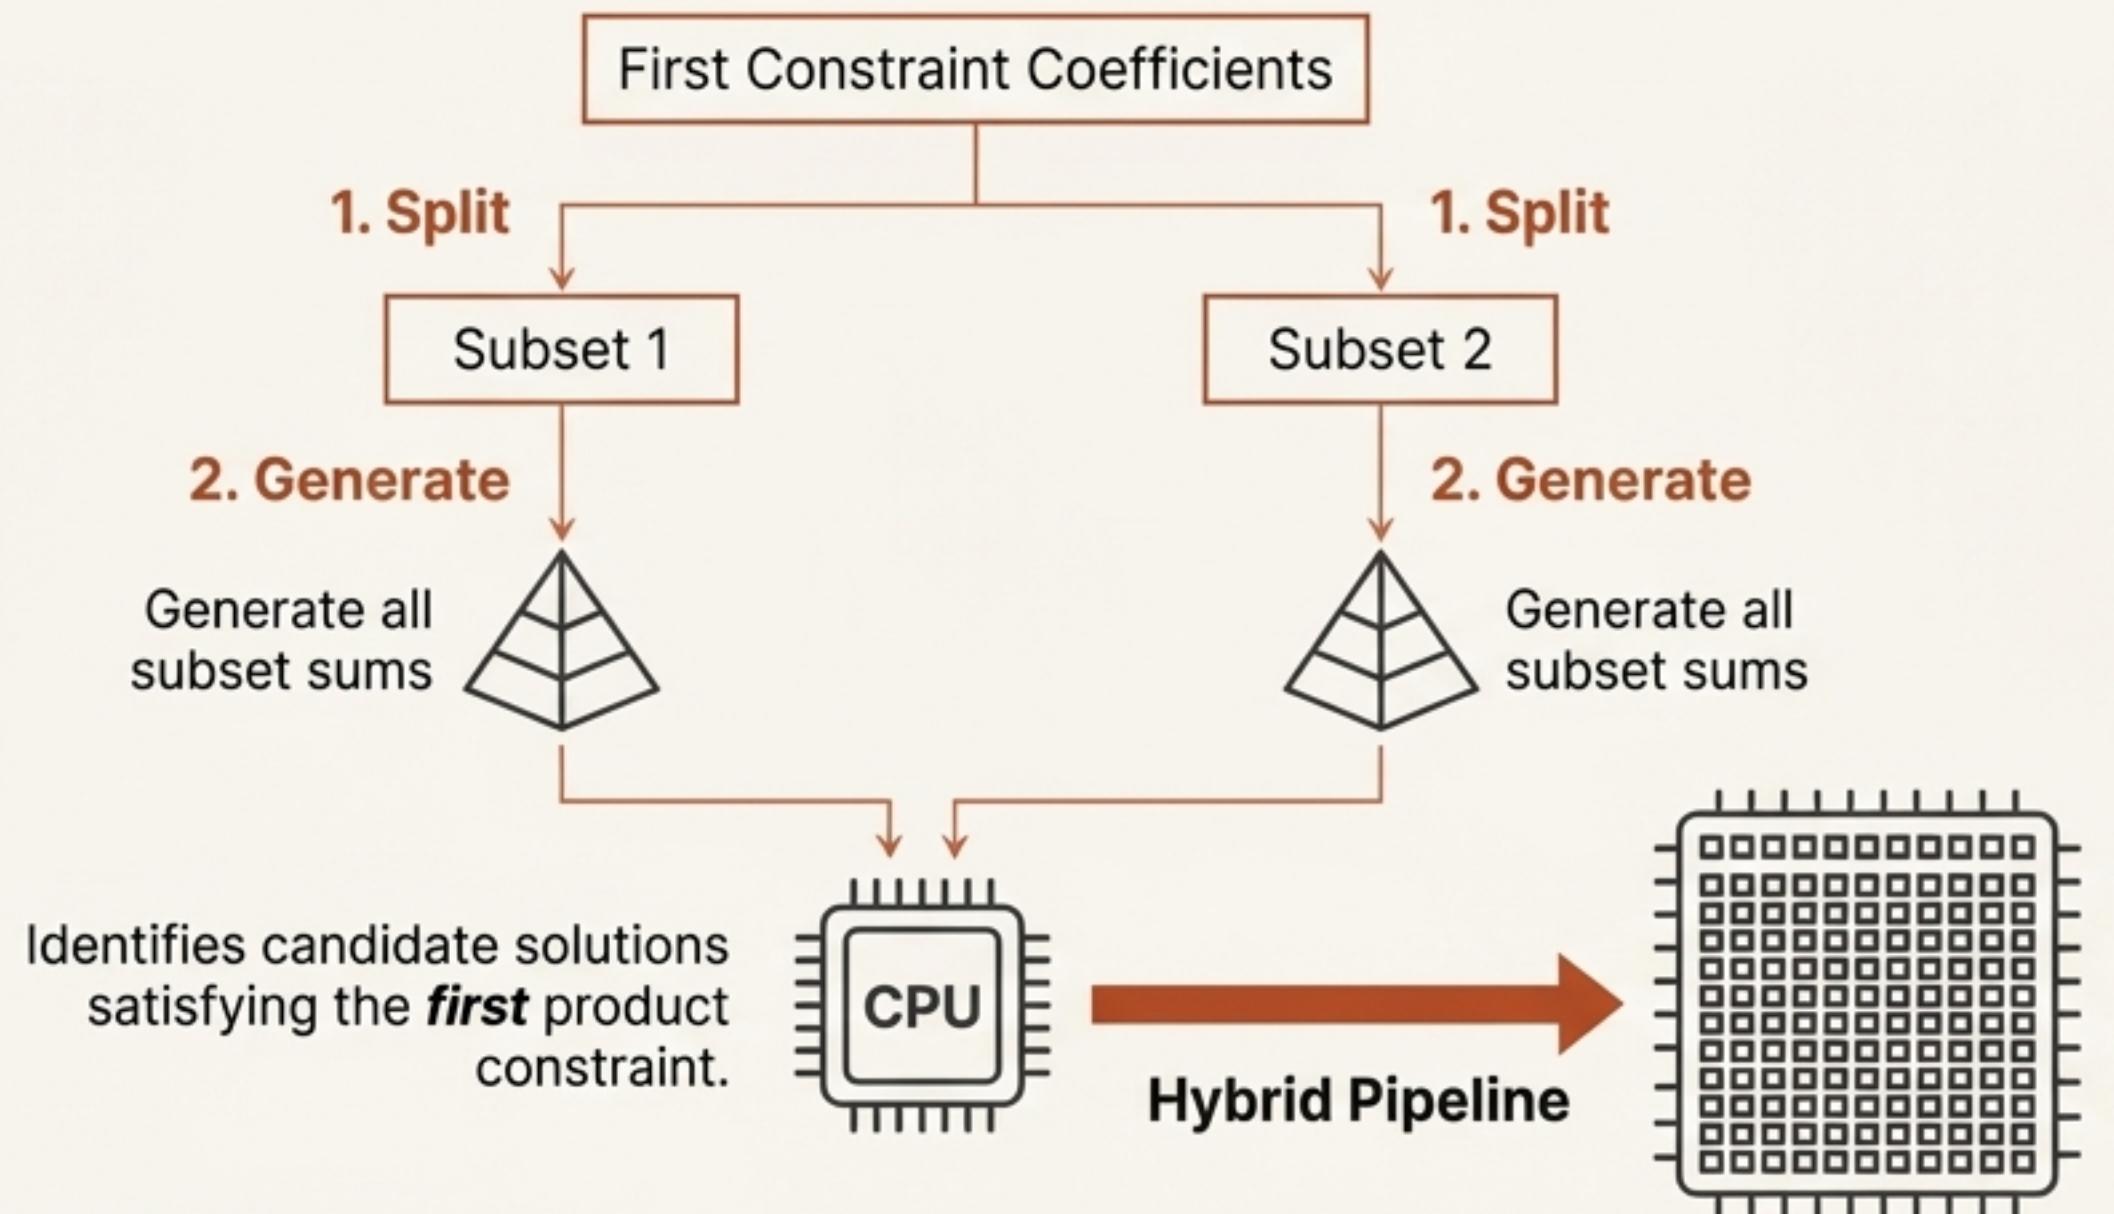
After reduction, it performs an exhaustive search of integer linear combinations of the basis vectors. This search is enhanced by using Holder's inequality for tighter bounds.

Advanced Search Strategy: Limited Discrepancy Search (LDS)

Recent versions use LDS as an alternative to depth-first search. LDS is highly effective for feasibility problems, as it prioritizes search paths that deviate minimally from a heuristic "ideal," often finding a solution much faster.

A Parallel Assault: The GPU-Accelerated Schroeppel-Shamir Algorithm

Instead of AHL-style lattice reduction, this method adapts a classic algorithm for the subset sum problem.



Performance Breakthroughs

This parallel approach has solved previously intractable instances:

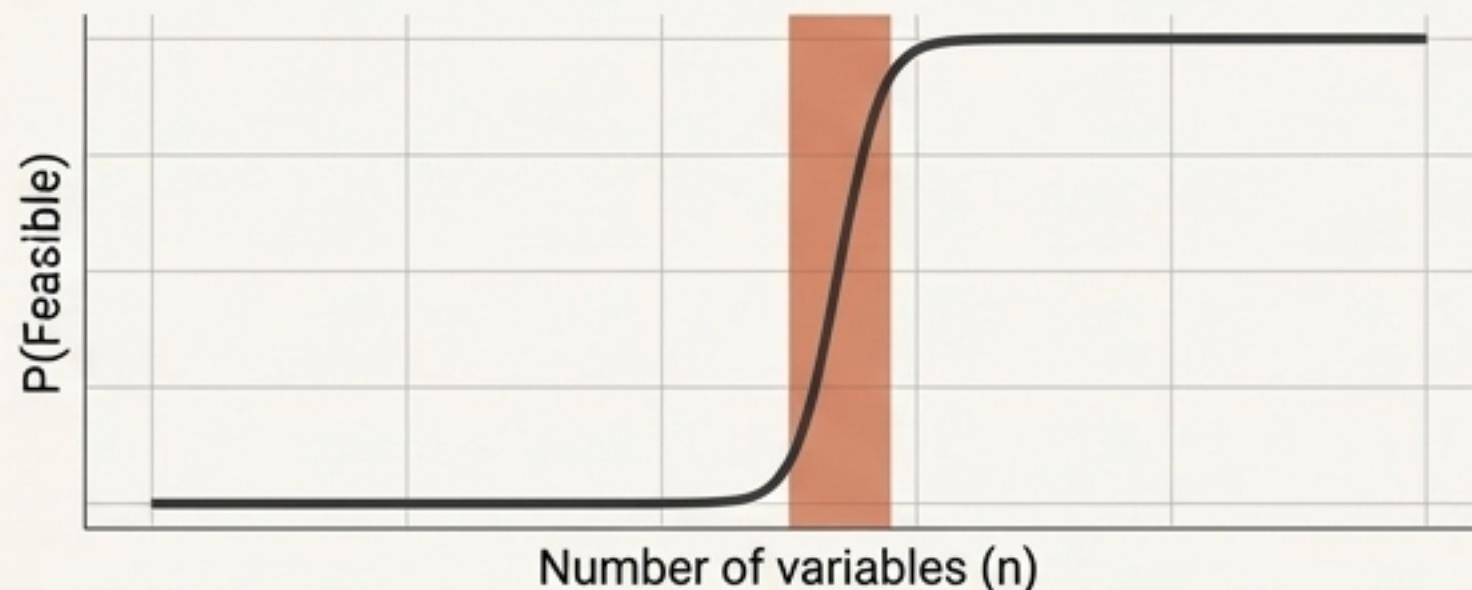
- **(9, 80)** in under **15 minutes**.
- **(10, 90)** in approximately **one day**.

GPU: Performs a massively parallel binary search and hash-based validation to check if these candidates satisfy the *remaining $m-1$* constraints.

Problem Difficulty is Governed by Phase Transitions and Coefficient Density

Beyond the ' $n = 10(m-1)$ ' Rule

Probabilistic analysis shows the probability of an instance being feasible undergoes a sharp phase transition as ' n ' increases.



The Role of Coefficient Range (D)

The range of demand coefficients ' a_{ij} ' significantly impacts solvability. Higher values of ' D ' increase problem "density," making solutions rarer and harder to find.

Demand Range (D)	Average Time (s)	Solutions Found
50	0.42	Multiple
100	12.80	Few
200	154.20	Rare

****Generating 'Hard' Infeasible Instances**:** For ' $m > 5$ ', the ' $10(m-1)$ ' rule actually produces many feasible instances. To create the hardest cases (which require a full proof of infeasibility), researchers now generate instances with slightly *fewer* variables than this threshold.

Data for ($m=6, n=50$) instances.

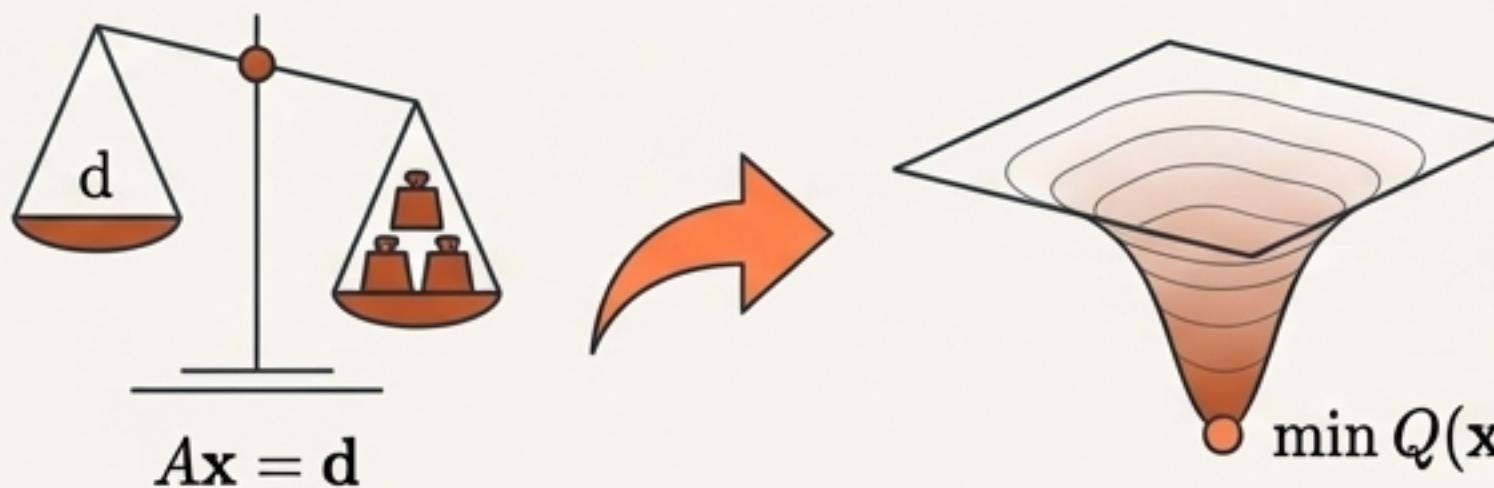
The Market Split Problem is Now a Primary Benchmark for Quantum Optimization

Why MSP?

Its characteristic difficulty at small variable counts makes it an ideal candidate for testing today's Noisy Intermediate-Scale Quantum (NISQ) devices.

The Transformation: From ILP to QUBO

To run on a quantum computer, the constrained linear system must be converted to a Quadratic Unconstrained Binary Optimization (QUBO) problem.



The QUBO Formulation

This is achieved by minimizing the squared violation of the constraints:

$$f(\mathbf{x}) = \sum_i \left(\sum_j A_{ij} x_j - d_i \right)^2$$

Expanding this yields the exact quadratic form $Q(\mathbf{x})$:

$$Q(\mathbf{x}) = \mathbf{x}^T (\mathbf{A}^T \mathbf{A}) \mathbf{x} - 2(\mathbf{A}^T \mathbf{d})^T \mathbf{x}$$

The global minimum of $Q(\mathbf{x})$ (with a value of zero) corresponds to a feasible solution of the original MSP. This allows the problem to be mapped directly to quantum annealers or solved via QAOA on gate-based hardware.

QOBLIB Features the MSP in its “Intractable Decathlon” for Quantum Hardware

Introducing QOBLIB: Established in 2025, the Quantum Optimization Benchmark Library (QOBLIB) provides a standardized set of challenges for quantum hardware.

Instance Curation: QOBLIB instances are categorized by the number of products m (3 to 15) and the demand range D (50, 100, 200).

A Key Guarantee: Every instance in the library is guaranteed to have at least one solution. This ensures that quantum solvers are tested on their ability to *find existing global minima*, not on their ability to *prove infeasibility*, which is currently beyond the scope of most quantum heuristics.



The Intractable Decathlon

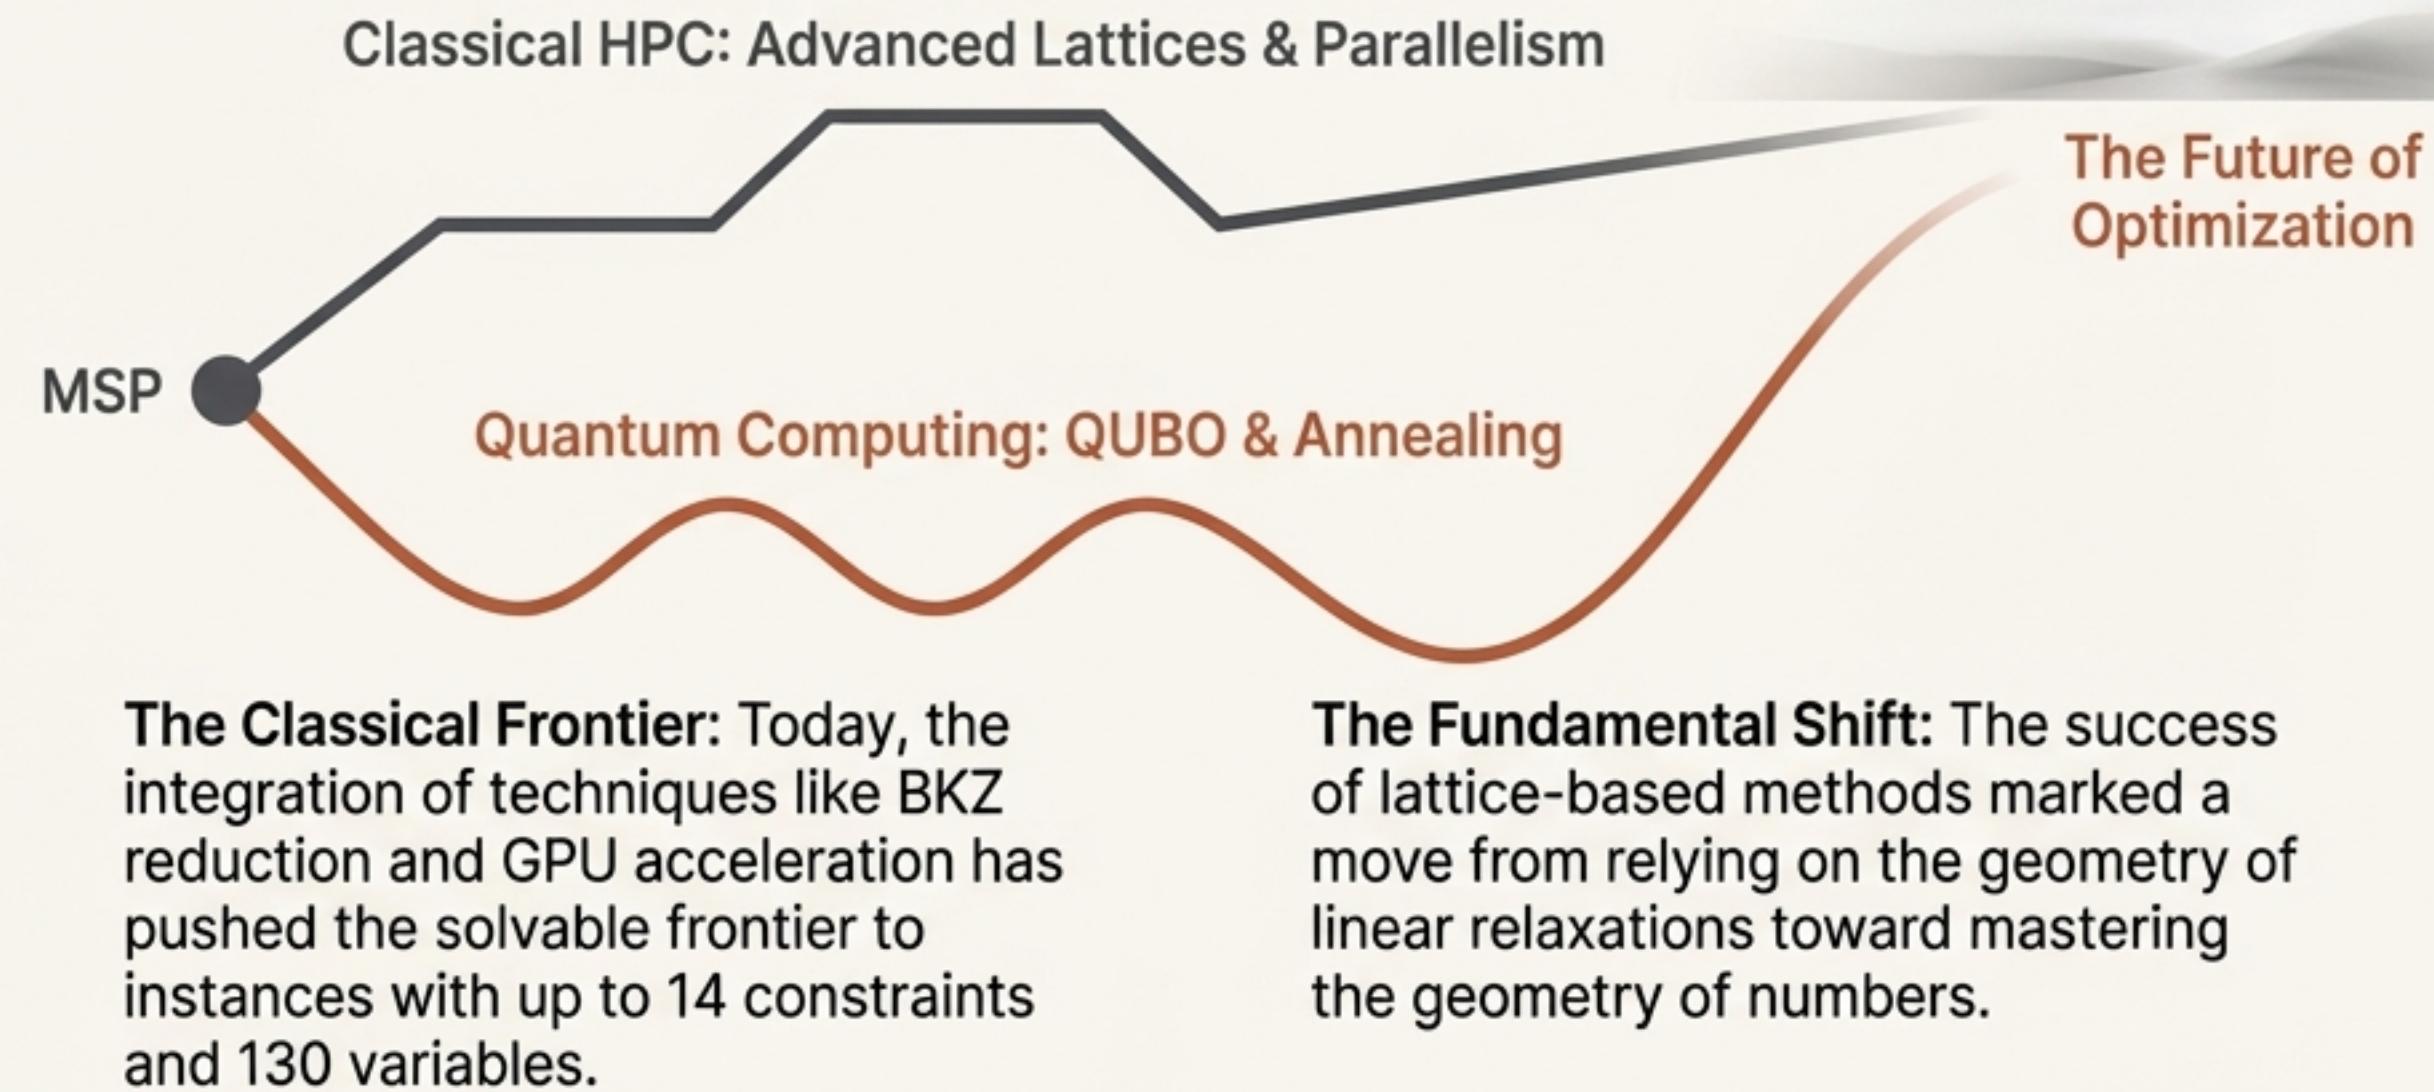
Number Partitioning
Graph Coloring
...

Market Split Problem
...

Maximum Independent Set

The Enduring Legacy of a Deceptively Simple Problem

The Journey Summarized: The MSP's evolution from a 1970s logistics problem to a 2020s quantum benchmark mirrors the broader evolution of how we tackle computational complexity.



A Definitive Measure of Ingenuity: The Market Split Problem will continue to serve as a crucible for algorithmic innovation, forcing constant refinement of our tools for navigating the intricate landscapes of discrete mathematics—in both classical and quantum computing.